

```

#ANN Without regularisation
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

# Load MNIST dataset
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()

# Normalize the pixel values
x_train, x_test = x_train / 255.0, x_test / 255.0

# Flatten the images (Since no hidden layers, direct input to output)
x_train = x_train.reshape(-1, 28*28)
x_test = x_test.reshape(-1, 28*28)

# Build the basic Perceptron model
model = keras.Sequential([
    layers.Dense(10, activation='softmax', input_shape=(28*28,))
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=32,
                    validation_data=(x_test, y_test))

# Evaluate model
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f"Test Accuracy: {test_acc:.4f}")

# Predict labels
y_pred = np.argmax(model.predict(x_test), axis=1)

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Plot confusion matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=range(10), yticklabels=range(10))
plt.xlabel("Predicted Label")

```

```

plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()

# Plot training vs validation accuracy/loss
plt.figure(figsize=(12, 4))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training vs Validation Accuracy')

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training vs Validation Loss')

plt.show()

```

Epoch 1/10

C:\Newanaconda\Lib\site-packages\keras\src\layers\core\dense.py:87:  
UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a  
layer. When using Sequential models, prefer using an `Input(shape)`  
object as the first layer in the model instead.

```

super().__init__(activity_regularizer=activity_regularizer,
**kwargs)

```

1875/1875 ————— 3s 1ms/step - accuracy: 0.8146 - loss:  
0.7197 - val\_accuracy: 0.9145 - val\_loss: 0.3086

Epoch 2/10

1875/1875 ————— 2s 902us/step - accuracy: 0.9137 -  
loss: 0.3116 - val\_accuracy: 0.9209 - val\_loss: 0.2826

Epoch 3/10

1875/1875 ————— 2s 813us/step - accuracy: 0.9202 -  
loss: 0.2843 - val\_accuracy: 0.9246 - val\_loss: 0.2705

Epoch 4/10

1875/1875 ————— 2s 881us/step - accuracy: 0.9229 -  
loss: 0.2776 - val\_accuracy: 0.9245 - val\_loss: 0.2709

Epoch 5/10

1875/1875 ————— 2s 886us/step - accuracy: 0.9281 -  
loss: 0.2603 - val\_accuracy: 0.9260 - val\_loss: 0.2654

Epoch 6/10  
 1875/1875 \_\_\_\_\_ 2s 897us/step - accuracy: 0.9269 -  
 loss: 0.2654 - val\_accuracy: 0.9273 - val\_loss: 0.2627

Epoch 7/10  
 1875/1875 \_\_\_\_\_ 2s 988us/step - accuracy: 0.9273 -  
 loss: 0.2619 - val\_accuracy: 0.9264 - val\_loss: 0.2646

Epoch 8/10  
 1875/1875 \_\_\_\_\_ 2s 988us/step - accuracy: 0.9297 -  
 loss: 0.2497 - val\_accuracy: 0.9251 - val\_loss: 0.2675


Epoch 9/10  
 1875/1875 \_\_\_\_\_ 2s 1ms/step - accuracy: 0.9317 - loss:  
 0.2491 - val\_accuracy: 0.9268 - val\_loss: 0.2627



Epoch 10/10  
 1875/1875 \_\_\_\_\_ 2s 1ms/step - accuracy: 0.9295 - loss:  
 0.2540 - val\_accuracy: 0.9272 - val\_loss: 0.2629

313/313 - 0s - 897us/step - accuracy: 0.9272 - loss: 0.2629

Test Accuracy: 0.9272

313/313 \_\_\_\_\_ 0s 961us/step





```
#With regularisation
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, regularizers
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

# Load MNIST dataset
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()

# Normalize the pixel values
x_train, x_test = x_train / 255.0, x_test / 255.0

# Flatten the images (Since no hidden layers, direct input to output)
x_train = x_train.reshape(-1, 28*28)
x_test = x_test.reshape(-1, 28*28)

# Build the Perceptron model with L1 & L2 regularization
model = keras.Sequential([
    layers.Dense(10, activation='softmax', input_shape=(28*28,),
        kernel_regularizer=regularizers.l1_l2(l1=0.01,
l2=0.01)) # L1 & L2 regularization
])

# Compile the model
model.compile(optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=32,
    validation_data=(x_test, y_test))
```

```

# Evaluate model
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f"Test Accuracy: {test_acc:.4f}")

# Predict labels
y_pred = np.argmax(model.predict(x_test), axis=1)

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Plot confusion matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=range(10), yticklabels=range(10))
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()

# Plot training vs validation accuracy/loss
plt.figure(figsize=(12, 4))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training vs Validation Accuracy')

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training vs Validation Loss')

plt.show()

```

```

C:\Newanaconda\Lib\site-packages\keras\src\layers\core\dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a
layer. When using Sequential models, prefer using an `Input(shape)`
object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer,
**kwargs)

```

```
Epoch 1/10
1875/1875 _____ 3s 1ms/step - accuracy: 0.7664 - loss:
2.1332 - val_accuracy: 0.8240 - val_loss: 1.6555
Epoch 2/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8151 - loss:
1.6729 - val_accuracy: 0.8105 - val_loss: 1.6519
Epoch 3/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8114 - loss:
1.6733 - val_accuracy: 0.8252 - val_loss: 1.6545
Epoch 4/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8143 - loss:
1.6725 - val_accuracy: 0.8176 - val_loss: 1.6457
Epoch 5/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8105 - loss:
1.6707 - val_accuracy: 0.8192 - val_loss: 1.6448
Epoch 6/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8119 - loss:
1.6702 - val_accuracy: 0.8126 - val_loss: 1.6503
Epoch 7/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8111 - loss:
1.6719 - val_accuracy: 0.8207 - val_loss: 1.6454
Epoch 8/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8109 - loss:
1.6705 - val_accuracy: 0.8255 - val_loss: 1.6510
Epoch 9/10
1875/1875 _____ 3s 1ms/step - accuracy: 0.8142 - loss:
1.6690 - val_accuracy: 0.8248 - val_loss: 1.6484
Epoch 10/10
1875/1875 _____ 2s 1ms/step - accuracy: 0.8159 - loss:
1.6697 - val_accuracy: 0.8176 - val_loss: 1.6504
313/313 - 0s - 861us/step - accuracy: 0.8176 - loss: 1.6504
Test Accuracy: 0.8176
313/313 _____ 0s 899us/step
```

